Green Adaptation of Real-Time Web Services for Industrial CPS within a Cloud Environment

M. Teresa Higuera-Toledano, José L. Risco-Martín, Patricia Arroba, and José L. Ayala

Abstract—Managing energy efficiency under timing constraints is an interesting and big challenge. This work proposes an accurate power model in data centers for time-constrained servers in Cloud computing. This model, as opposed to previous approaches, does not only consider the workload assigned to the processing element, but also incorporates the need of considering the static power consumption and, even more interestingly, its dependency with temperature. The proposed model has been used in a multi-objective optimization environment in which the Dynamic Voltage and Frequency Scaling (DVFS) and workload assignment have been efficiently optimized.

Index Terms—Adaptive Systems, Cyber Physical Systems, Cloud Computing, Real-Time Systems, Energy efficiency, Industrial-based Services, Multi-Objective Optimization, Parallel Computing.

I. INTRODUCTION

B OTH *Cyber Physical Systems* (CPSs) and *Cyber Physical Society* [1] combine computing and networking power with physical components, enabling innovation in a wide range of domains related to future-generation sensor networks (e.g., robotics, avionics, transportation, manufacturing processes, energy, smart homes and vehicles, medical implants, healthcare, etc). The design and implementation of CPS involve the consideration of multiple aspects like energy and tight real-time constraints. Because of that, real-time scheduling for CPS brings new research issues in the scope of real-time systems [2].

Managing energy efficiency under timing constraints is a big challenge. Most modern micro-controllers already provide support for various energy saving modes (e.g., Intel Xeon and AMD Opteron). A common way of reducing dynamic power is to use the technique called *Dynamic Voltage and Frequency Scaling* (DVFS), which changes the processor voltage and the clock frequency simultaneously, reducing the energy consumption. Decreasing the processor voltage and frequency will slow down the performance of the processor. If the execution performance is not a hard constraint, then, decreasing both processor voltage and frequency allows to reduce the dynamic power consumption of the processor.

M. Teresa Higuera-Toledano, José L. Risco-Martín, and José L. Ayala are with the Department of Computer Architecture and Automation, Complutense University of Madrid, C/Prof. José García Santesmases 9, 28040 Madrid, Spain, email: mthiguer,jlrisco,jlayalar@ucm.es

Patricia Arroba is with the Department of Electronic Engineering, Technical University of Madrid, Avda. Complutense 30, 28040, Madrid, Spain, email: parroba@die.upm.es

Copyright (c) 2009 IEEE. Personal use of this material is permitted. However, permission to use this material for any other purposes must be obtained from the IEEE by sending a request to pubs-permissions@ieee.org.

Nowadays, new embedded devices are collaborating in distributed environments. In this new scenario, tasks and resources are widely distributed and then, real-time applications become more complex and more relevant. A cloud datacenter usually contains a large group of servers connected through the Internet, and a scheduler has to make an efficiently use of the resources of the cloud to execute jobs. Since many applications require *Quality of Service* (QoS), power consumption in data centers must be minimized, satisfying the *Service Level Agreement* (SLA) constraints. Consequently, novel approaches that base their optimizations on accurate power models must be devised, performing an optimized setting of the parameters of the server (frequency, voltage, workload allocation, etc) while accomplishing with time requeriments and a wide range of real-time constraints.

DVFS-based solutions for distributed real-time environments identify two main dimensions of the problem: (i) task-to-Central-Processing-Unit (CPU) allocation and (ii) run-time voltage scaling on individual CPUs. In CPS, physical factors (e.g., the network topology of CPS may dynamically change due to physical environments) are not entirely predictable and may lead to problems such as missed task deadlines, that can impact dramatically on economic loss for individuals or for the industry. Moreover, a critical task deadline missed could trigger a disaster (e.g., humans life loss, natural disasters, or huge economic loss).

In this paper, we propose a method for solving such CPS problems by introducing new adaptive real-time scheduling algorithms in distributed computing infrastructures that also consider energy efficiency. This scheme requires to know *a priori* the processing and timing constraints of the set of tasks, and must be supported by reservation-based real-time operating systems.

The remainder of this paper is organized as follows: after a brief summary of the previous works in this field (Section II), a real-time scheduling algorithm for CPS is sketched (Section III). Following, the devised power model is presented (Section IV), and the optimization of the algorithm developed is profusely described (Section V). Experimental results can be found in Section VI. Finally, some conclusions are drawn (Section VII).

II. RELATED WORK

The energy-efficient scheduling problem in real-time systems consists in minimizing the energy consumption while ensuring that all the real-time tasks meet their deadlines. The work presented in [3] is based on the observation that a

significant percentage of time spent in idle mode is due to the accumulation of small gaps between tasks. Whether the gap between the activation of two periodic tasks is less than transition-time from idle to deep-sleep, the processor is not able to transition to the deep-sleep state even though there is no useful work to be done, and continues in the idle energy state all the time.

There are extensive research works on energy-aware realtime scheduling by using DVFS (e.g., [4]). Different works using this technique within a real-time context, considered an offline scheduling algorithm and a set of a periodic jobs on an ideal processor. Each job is characterized by its release time, deadline, and execution CPU cycles, and all jobs have the same power consumption function.

Several papers have also proposed DVFS-based solutions for real-time multi-processor systems. As the complexity of CPS increases, *Chip Multicore Processors* (CMP) and parallel tasks scheduled in a real-time way are needed. The fact that the processing cores share a common voltage level makes the CMP energy-efficiency problem different from multi-processor platforms. The work presented in [5] provides a simple, elegant and effective solution on energy-efficient real-time scheduling on CMP. This solution addresses fixed-priority scheduling of periodic real-time tasks having a deadline equal to their period. Note that this problem is NP-hard.

The load balancing in CMP is particularly important because the main contributor to the overall energy consumption in the system is the core with the maximum load. This fact is given by the global voltage/frequency constraint. Considering a CMP system with a workload perfectly balanced across the processors, the *Earliest Deadline First* (EDF) scheduling minimizes the total energy consumption. This is not the case of *Rate Monotonic Scheduling* (RMS) where load-balancing does not always result in lowering energy consumption [5].

In mixed-criticality systems, varying degrees of assurance must be provided to functionalities of varying importances. As shown in [6] there is a conflict between safety and energy minimization because critical tasks must meet their deadlines even whether exceeding their expected *Worts Case Execution Time* (WCET). This work integrates continuous DVFS with the EDF with *Virtual Deadlines* (EDF-VD) scheduling for mixed-criticality systems [7] and shows that speeding up the system to handle overrun is beneficial for minimizing the expected energy consumption of the system.

Generally, large-scale distributed applications require realtime responses to meet soft deadlines. Hence, the middleware coordinates resource allocation in order to provide services accomplishing with SLA requirements. In [8], we can find a scheduling algorithm based on DVFS for clusters, which develops a green SLA-based mechanism to reduce energy consumption by increasing the scheduling makespans. In [9], we can find an energy-aware resource allocation for Cloud computing with negotiated QoS. However, similarly to the solution presented in [8], this method sacrifices system performance.

The work presented in [10] proposes a priority-based scheduler, which satisfies the minimum resource requirement of a job by selecting a *Virtual Machine* (VM) according to both

the SLA level and the W_i parameter that is described as $W_i = P_i \times R_i$, where P_i is the unit power cost of VM_i , and R_i defines the resources used by the VM_i .

The location of nodes in CPS affects the effective release time and deadline of real-time tasks, which may be different depending on the node location and the migration delay time among the network nodes. Because of that, traditional real-time scheduling algorithms have to be modified to include the location node and the spatial factors. The work presented in [11] proposes a CPS scheduling algorithm, where the servicing node (i.e., the CPU) needs to move to serviced (i.e., the executed Job) node for real-time services.

The power modeling technique proposed in [12] is most relevant for us. A correlation between the total system's power consumption and the component utilization is observed, defining a four-dimensional linear weighted power model for the total power consumed (i.e., $P = c_0 + c_1 P_{CPU} + c_2 P_{cache} + c_3 P_{DRAM} + c_4 P_{disk}$). Our work follows a similar approach but also incorporates the contribution of the static power consumption, its dependency with temperature, and the effect of applying DVFS techniques.

Static power consumption has a high impact on energy, due to the temperature-dependent leakage currents. In this manner, novel optimizations may be devised by quantitatively understanding the power-thermal trade-offs of a system, thus developing analytical models.

Finally, Rafique et al. [13] makes a description of the complexity of the power management and allocation challenge. Authors demonstrate that achieving an optimal allocation depends on many factors as the server's maximum and minimum frequencies, the job's arrival rate, and consequently, the relationship between power and frequency. They conduct a set of experiments that provides significant savings in terms of energy in both homogeneous and heterogeneous clusters. However, our work presented in this paper outperforms these savings by exploiting a multi-objective optimization strategy to help to minimize the servers' power for time-constrained Cloud applications.

III. THE INDUSTRIAL SERVICES EXECUTION MODEL

CPS comprise a large number of sensors and actuators, and computing units that exchange different types of data, some of these interactions have real-time constraints. Real-time system abstraction and hybrid system modeling and control are among the CPS research challenges. The hybrid system model of CPS requires the design and integration of both the physical and computational (i.e., cyber) elements. While physical elements behave in continuous real-time, computational elements change according to discrete logic. This fact requires to merge continuous-time based systems with event-triggered logical systems, and also we must address the dimensional scale (i.e., from on-chip level to the cloud). Moreover, the interaction with physical world introduces uncertainty in CPS because of randomness in the environment, errors in physical devices, and security attacks.

Control and scheduling co-design is a well-known area in the embedded real-time systems' community. However, since CPS are typically networked control systems, the tradeoff between the effects of the network must be included in the real-time schedulability, that results in a non-periodic control approach. In this work, we study how to guarantee the overall system stability with minimum computational resource and power usage. System properties and requirements (e.g., the control laws, real-time and power constraints) must be captured and supported by data abstractions encapsulated in components.

A. Task characterization

Typically CPS's are composed of hard real-time tasks and feedback control tasks. Whereas real-time tasks present time constraints (i.e., deadlines) that must always be satisfied, feedback control tasks are characterized by their *Quality of Control* (QoC), which needs to be optimized. A typical approach to the above scheduling problem is to translate the QoC requirements into time constraints and then, to apply traditional real-time scheduling techniques [14]. Real-time systems are structured as a set of schedulable tasks, where parameters used for the scheduling (e.g., execution time, deadline, or period) are a priori known and clearly defined. However, this solution is very conservative and consequently it is not efficient for CPS.

An alternative solution is the given in [15], that deals with this problem using a *multi-layered* scheme based on mixed-critical real-time systems: (i) for real-time tasks it uses triggering patterns (i.e., uses arrival curves), which allow a more general characterization regarding the classical real-time task models (i.e., *periodic* or *sporadic*), and (ii) for control tasks, it is based on three QoC-oriented metrics. Mixed-critical real-time systems literature focuses on tasks with different criticality levels and certification issues¹, providing heterogeneous timing guarantees for tasks of varying criticality levels.

As an example, in the *Unmanned Aerial Vehicles* (UAVs), functionalities can be categorized as safety-critical tasks (e.g., like flight control and trajectory computation) or mission-critical tasks (e.g., object tracking for surveillance purposes). Note that the system is still safe although mission-critical functionalities can be lost. This makes the design parameters for safety-critical tasks (e.g., WCET) much more pessimistic than those for mission-critical tasks. However, in CPS, tasks are not characterized by criticality levels, but by their criticality types.

There has been considerable research on schedule synthesis for control applications. However, these works are particularly centered on control/scheduling co-design for optimized QoC, and only deal with control tasks. On the other hand, CPS focus on mixed task sets comprising of feedback control tasks and hard real-time tasks, which requires a joint schedule synthesis.

B. The task model

In CPS, tasks may be classified according to their criticality types (e.g., deadline-critical real-time tasks and QoC-critical

¹When there are tasks with different safety requirements into the same real-time platform, it is called mixed-criticality system.

feedback control tasks). While the system must satisfy always the deadlines of real-time tasks, particularly for those that are critical, only the QoC parameters for control tasks need to be optimized. In order to do that, we require stochastic hybrid systems to identify the interaction between continuous dynamical physical models and discrete state machines, and the CPS architecture must follow the new paradigm "globally virtual, locally physical".

We consider a set of independent tasks, (i.e., Σ) which are executed remotely in a set of physical servers m. We define our real-time problem as a pair $P=(\Sigma,S)$ where S is a scheduling solution and $\Sigma=\tau_1,...,\tau_n$ is a set of n tasks with different timing characteristics (i.e., strict, flexible, and firm) as shows Figure 1.

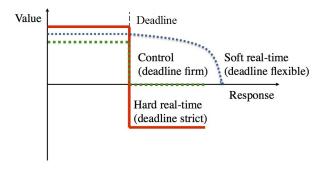


Fig. 1. An overrun in response time (i.e., a deadline miss) has a different value function depending on its possible consequences

Each task τ_i is a possibly infinite sequence of jobs (i.e., demands for processing time), each one with an associated deadline. Jobs of the same task must be executed sequentially and in First-In-First-Out (FIFO) order. If the timing characteristics of the task τ_i are soft or firm, the jobs may be not identical

CPS requires jointly scheduling hard real-time, soft real-time or best-effort, and control-feedback tasks. Due to the stringent stability requirements, we classify control tasks as firm deadline. While a hard deadline cannot be missed, soft deadlines may be occasionally missed and it does not harm the system safety. Similarly, firm deadlines can be missed but there is an upper limit on the number of misses within a given time interval. However, as we aim to optimize the QoC, we must minimize the number of deadline misses to avoid QoC degradation. The characterization of each type of task is fundamentally different as follows.

1) Hard real-time tasks: A real-time system is considered hard if an overrun in a task response time leads to potential loss of life and/or big financial damage. The system is considered to be safety critical or high integrity, and is often mission critical. We consider a real-time task as a tuple $\tau_i = (R_i, C_i, T_i, D_i)$ where:

- R_i is the first release time of the task (i.e. the *phase* of the task).
- C_i is the WCET,
- T_i is the activation period (i.e., minimum inter-release time), and

We compute the CPU utilization factor of τ_i as $U_i = \frac{C_i}{T_i}$.

- 2) Soft real-time tasks: For soft real-time tasks, deadline overruns are tolerable but not desired (i.e., there are not catastrophic consequences of missing one or more deadlines). There is a cost function associated with these systems, which is often related to QoS. Hence, we consider a stochastic task model based on the one presented in [16]. Then, we represent each soft-real-time task using a tuple $\tau_i = (r_i, s_i, a_i, d_i)$ where:
 - r_i is the release time of the task,
 - is the service time, which follows an exponential S_i distribution of average μ^{-1} (i.e., μ is the number of serviced jobs of τ_i per unit time),
 - is the arrival time; tasks arrive according to a renewal a_i process with exponential distribution of average λ^{-1} ,
 - is the absolute deadline; the relative deadline is $D_i =$ d_i $d_i - a_i$, D_i distributed on [0, D].

We compute the response time of τ_i as $\rho_i = c_i - a_i$, where c_i is the completion time (i.e., $c_i = a_i + s_i$). The average CPU utilization factor is given by $\Upsilon_i = \frac{\mu_i}{\lambda_i}$.

3) Feedback control tasks: For a firm real-time task the computation is obsolete whether the job is not finished on time. In this case, the cost function may be interpreted as loss of value associated to QoC. This is the case of the feedback control task in CPS. For this kind of task we can consider $D_i > T_i$) to guarantee that the controlled physical tasks are still stable in the worst case scenario. However, this sacrifices the system performance and also may result unstable under physical perturbations.

In most cases, feedback control systems become unstable with too many missed control cycles. Therefore, a critical question is how to determine T_i to ensure both schedulability and robustness of the physical system. Considering a simple proportional-gain feedback controller, which is fixed for each control task, in order to determine T_i , we can find the minimum $T_i \in (T_1, T_2, \dots, T_n)$ under the following constraints:

$$0 \le \sum_{i} \frac{C_{i}}{T_{i}} \le p$$

$$C_{i} \le T_{i} \le D_{i}$$

$$(1)$$

$$C_i < T_i < D_i \tag{2}$$

where p < 1 is a priori known. However, some controller parameters may need to be adjusted when the task period is changed. Alternatively, we can use a multiple-versions approach or a predictive model with a quadratic optimization computed iteratively for each job. However, very often, probabilistic guarantees are sufficient (e.g., t out of k deadlines have to be met).

Permitting skips in periodic tasks increases the system flexibility [17]. The maximum number of skipped jobs for each task can be controlled by a specific parameter S_i associated with the task, which gives the minimum distance between two consecutive jobs skips (e.g., if $(S_i = 3)$ the task can skip one job every three). This parameter can be considered as a QoC metric (i.e., the higher S, the better QoC).

When $S_i = \infty$ no skips are allowed, meaning that τ_i is a real-time hard periodic task. We then consider a control task as a tuple $\tau_i = (R_i, C_i, T_i, D_i, S_i)$ where $T_i = D_i$.

C. The parallel scheduling

For each of the above described tasks $\tau_i \in \Sigma$, we consider a set of independent subtasks $\tau_i = \tau_{i,1}, ..., \tau_{i,q}$, where $\tau_{i,j}$ denotes the subtasks j of task τ_i . Therefore, $e_i \geq 0$, is the energy consumption rate of the task τ_i per time unit:

$$e_i = \bigcup_{j=1}^{|\tau_i|} e_{i,j} \tag{3}$$

The scheduling allocates each $\tau_{i,j}$ subtask in a set of m physical servers, taken into account the critical timing characteristics of each task τ_i and the minimal energy consumption of the task set Σ .

The performance criteria generally used in systems when the model task does not have explicit deadlines, is to minimize the task delay (i.e., the response time of all tasks). However, when there are explicit deadlines, we must ensure that critical tasks fulfill their deadline and minimize the fraction of non-critical tasks that do not meet their timing requirements.

We can consider lateness constraints of the form $\alpha(x) \leq \beta$, where $\alpha(x)$ is the fraction of jobs that miss their deadline by more than x time units. Here, missing a deadline by x time units is considered as a failure.

- For firm deadlines, we require that $\alpha(0) < \beta$ (i.e., the fraction of tasks missing their deadliness were limited to β). Note that this has a different meaning for the S parameter, which is the minimal distance between the consecutive misses of the task τ_i . Hence, we consider a τ_i missing whether one or more subtasks $\tau_{i,j}$ of a job miss the deadline.
- For hard real-time tasks, we establish $\alpha(0) \leq 0$ (i.e., we do not tolerate any deadline missed), while for each control task τ_i , $\alpha(0) \leq \frac{S_i - 1}{S_i}$.
- For soft real-time tasks, we generalize, $\alpha(x_i) \leq \beta_i$, for a set of time values $x_1, ..., x_p$ and constraint specifications $\beta_1, ..., \beta_p$, where $1 \le i \le p$, which allows to take into account the stochastic nature of task arrivals and service time of soft real-time tasks.

IV. POWER AND ENERGY MODEL

Traditionally in electronic systems, dynamic consumption has been the major contributor to the power budget. In contrast, when scaling technology below 100nm, static consumption reaches the 30-59% of the total power, thus becoming much more significant [18]. Moreover, the exponential impact of temperature on leakage currents intensifies this effect. Thus, modeling leakage will allow the exploitation of the trade-offs between leakage and temperature at the server level when taking decisions on resource configuration and selection.

Therefore, the impact of static consumption must be considered, taking into account its correlation with temperature. This section presents our leakage-aware static power model. We validate this model using real data gathered from real machines of our case study (e.g., Intel Xeon and AMD Opteron).

A. Leakage power

Equation (4) shows the impact of leakage on the currents in a MOS device. Rabaey demonstrates in his work that, when $V_{DS} > 100 mV$, the second exponential may be considered negligible [19]. Consequently, the previous equation may be revised as in (5), also regrouping technological parameters together obtaining the formula presented in equation (6).

$$I_{leak} = I_s \cdot e^{\frac{V_{GS} - V_{TH}}{nkT/q}} \cdot \left(1 - e^{\frac{V_{ds}}{kT/q}}\right) \tag{4}$$

$$I_{leak} = I_s \cdot e^{\frac{V_{GS} - V_{TH}}{nkT/q}}$$

$$I_{leak} = B \cdot T^2 \cdot e^{\frac{V_{GS} - V_{TH}}{nkT/q}}$$
(6)

$$I_{leak} = B \cdot T^2 \cdot e^{\frac{v_{GS} - v_{TH}}{nkT/q}} \tag{6}$$

The leakage power consumption for the physical machine $m \in \{1,\ldots,M\}$ presented in Equation (8) can be inferred from the expression in (7). Then, the expansion of the mathematical expression in its Taylor 3rd order series provides Equation (9), where B_m , C_m and D_m represent the technological constants of the server.

$$P_{leak,m} = I_{leak,m} \cdot V_{DD,m} \tag{7}$$

$$P_{leak,m} = I_{leak,m} \cdot V_{DD,m}$$

$$P_{leak,m} = B_m \cdot T_m^2 \cdot e^{\frac{V_{GS} - V_{TH}}{nkT/q}}$$
(8)

$$P_{leak,m} = B_m \cdot T_m^2 \cdot V_{DD,m} + C_m \cdot T_h \cdot V_{DD,m}^2 + D_m \cdot V_{DD,m}^3$$
 (9)

B. Dealing with DVFS

The main contributors to energy consumption in nowadays servers are CPU and memory devices. Despite DVFS is easily found in CPUs, there are still few memories with these capabilities. However, memory consumption in some cases (memoryintensive applications) is very significant compared to the CPU consumption and, because of this, it was considered important enough to be studied independently.

Equation 10 provides the consumption of a physical server that has $k \in \{1 \dots K\}$ DVFS modes, while memory remains at a constant voltage. This expression takes into account the impact of temperature on the static power contribution. We define E_m as the contribution of other server resources operating at constant values of frequency and voltage.

$$\begin{array}{lcl} P_{leak,mk} & = & B_m \cdot T_{CPU,m}^2 \cdot V_{DD,mk} \\ & + & C_m \cdot T_{CPU,m} \cdot V_{DD,mk}^2 + D_m \cdot V_{DD,mk}^3 \\ & + & E_m + G_m \cdot T_{MEM,m}^2 + H_m \cdot T_{MEM,m} \end{array} \label{eq:pleak}$$

In order to measure temperature-dependent leakage we must understand also the dynamic contribution of the server's power consumption. To maintain constant conditions, we use lookbusy 2, which is a synthetic application that stresses the CPU during specifics periods of time. Lookbusy is able to stress, not only the cores but also the hardware threads of the CPU at a precise utilization, having no impact on memory or disk devices. Synthetic workloads help us to maintain the utilization rate constant (in terms of instructions per cycle), thus revealing the leakage contribution due to temperature variations. The formulation of the dynamic power consumption is shown in Equation 11.

$$P_{CPU,dyn,imk} = A_m \cdot V_{DD,mk}^2 \cdot f_{mk} \cdot u_{CPU,imk}$$
 (11)

where A_m defines the technological constant of the physical machine m and f_{mk} and $V_{DD,mk}$ are respectively the frequency and the supply voltage at the k DVFS mode of the CPU. $u_{CPU,imk}$ represents the CPU utilization and it is correlated with the number of CPU cycles.

C. Energy model

So far, the power model is derived as in (12).

$$P_{tot,mk} = A_m \cdot V_{DD,mk}^2 \cdot f_{mk} \cdot \sum_i u_{CPU,imk}$$

$$+ B_m \cdot T_{CPU,m}^2 \cdot V_{DD,mk}$$

$$+ C_m \cdot T_{CPU,m} \cdot V_{DD,mk}^2$$

$$+ D_m \cdot V_{DD,mk}^3 + E_m$$
(12)

The corresponding energy model can be easily obtained taking into account that $E = P \times t$, being P the power model in (12) and t, the execution time. Thus, the total energy consumed per host is described as the summation of the following equations:

$$E_{CPU,dyn,mk} = A_m \cdot V_{DD,m}^2 \cdot CPI \cdot \sum_{i} u_{CPU,imk} \cdot n_{CPU,imk}$$
(13)
$$E_{leak,mk} = [B_m \cdot T_{CPU,m}^2 \cdot V_{DD,m} + C_m \cdot T_{CPU,m} \cdot V_{DD,m}^2 + E_m + G_m \cdot T_{MEM,m}^2 + H_m \cdot T_{MEM,m}$$

$$[] \cdot \frac{CPI}{f_{mk}} \cdot \sum_{i} n_{CPU,imk}$$
(14)

where

- CPI is the number of cycles per instruction
- $n_{CPU,imk}$ is the number of CPU instructions of each task i assigned to be executed in a specific server m and DVFS mode k.

The summation of both the instructions to execute and the resources used by the workload hosted on the server are needed in order to get the execution time of all tasks executed in parallel considering the resources offered by each server, as seen in (14).

$$E_{tot} = \sum_{mk} \left(E_{CPU,dyn,mk} + E_{leak,mk} \right) \tag{15}$$

V. MULTI-OBJECTIVE OPTIMIZATION ALGORITHM

In this work, we aim for a workload allocation in a cloud that allows to optimize energy consumption. In addition, the benefits offered by virtualization are exploited, allowing to

²http://www.devin.com/lookbusy/

allocate the tasks in a more versatile way. The proposed system is defined as a cluster of machines of a cloud facility.

The proposed solution considers server heterogeneity, so the technological parameters will vary from one architecture to another, resulting in a different energy consumption. Since the resultant power model is non-linear and there exists a large set of constraints, the problem is tackled as a multi-objective optimization:

Minimize

$$\mathbf{y} = \mathbf{f}(\mathbf{x}) = [\lambda, (1+\lambda) \cdot E_{tot}(\mathbf{x})]$$

Subject to
 $\mathbf{x} = (x_1, x_2, \dots, x_n) \in \mathbf{X}$ (16)

where \mathbf{x} is the vector of n decision variables, \mathbf{f} is the vector of 2 objectives function, λ is the number of constraints not satisfied, E_{tot} is the total energy, and \mathbf{X} is the feasible region in the decision space. Using λ as shown in Equation 16, unfeasible solutions are also allowed, but only when no other alternatives are found. In this particular case, E_{tot} is measured using (15), whereas λ is computed as a penalization over the control and soft tasks that are delivered after the deadline (see Figure 1).

Using this formulation, we are able to obtain optimal energy savings, realistic with the current technology. To provide an efficient assignment in data centers it is necessary to consider both the energy consumption and the resource needs of each task of the workload.

A task τ_i can be split in different subtasks $\tau_{i,j}$ in order to achieve energy savings. Therefore, a task τ_i can be executed using a specific amount of resources of one or more servers defined by $u_{CPU,imk}$. The utilization percentage of the resources assigned to a task determines its execution time (i.e., C_i or s_i). In summary, the proposed multi-objective formulation, once solved, decides the following aspects:

- Operating server set, indicating which hosts are active according to the operating conditions of each physical machine.
- Best assignment for the various tasks of the workload, distributing each CPU instruction and memory requirements according to the minimum energy consumption of the applications in the computing infrastructure. For control tasks, S=2 must be fulfilled. However, a penalty is added to λ when one control task is aborted, even when S is being satisfied.
- Percentage of resources used by every task in each host where it is allocated, achieving best energy consumption.

A. The solver

Evolutionary algorithms have been used to run the proposed multi-objective formulation. In this work, we use the Non-dominated Sorting Genetic Algorithm II (NSGA-II) [20], which has become a standard approach to solve this kind of problems [21]. The chromosome encoding is shown in Figure 2.

In this case, each gene represents a decision variable. Because many decision variables are integer, the chromosome uses

$DVFS_1 \mid \cdots \mid DVFS_M$	$n_{CPU,11}$ · ·	$n_{CPU,NM}$
----------------------------------	------------------	--------------

Fig. 2. Chromosome encoding

integer encoding. Decision variables like $n_{CPU,imk}$ are scaled to the integer interval $0 \le n_{CPU,imk} \le 100$, and transformed to its real value (i.e., multiplying the percentage by the total number of instructions in the multi-objective function for evaluation).

NSGA-II is always executed with an initial random population of 100 chromosomes. After that, the algorithm evolves the population applying (1) the NSGA-II standard tournament operator, (2) a single point crossover operator with probability of 0.9 as recommended in [20], (3) a integer flip mutation operator (with probability of 1/number of decision variables as also recommended in [20], and (4) the multi-objective evaluation. Steps (1) to (4) are applied for a variable number of iterations or generations, which depend on the time that the parameter λ becomes 0 (usually 25000 iterations have been enough).

VI. RESULTS

Tests have been conducted gathering real data from a Fujitsu RX300 S6 server based on an Intel Xeon E5620 processor and a SunFire V20z Dual Core AMD Opteron 270, both operating at the set of frequencies f_{mi} given in Table I. Total power consumption and CPU temperature have been collected via the *Intelligent Platform Management Interface* (IPMI) during the execution of *lookbusy* at different utilization levels ranging from 0% to 100%, where a 65% of these levels were used to fit the energy model and the remaining 35% for validation. We used MATLAB to fit our data, obtaining the constants and validation errors shown in Table II.

TABLE I INTEL XEON E5620 AND SUNFIRE V20z DUAL CORE AMD OPTERON 270 FREQUENCIES

Platform	f_{m1}	f_{m2}	f_{m3}	f_{m4}	f_{m5}	f_{m6}
Intel (GHz)	1.73	1.86	2.13	2.26	2.39	2.40
AMD (GHz)	1.0	1.8	2.0			

The efficiency of the power supplies affects the calculation of these constants for different temperatures. In consequence, negative constants appear due to the fact that only CPU and memory have been characterized in this work because of their dominant contribution. In order to adapt the problem to more generic Cloud computing environments, our model constants can be calculated for data obtained during the execution of the workload in virtual machines. In that experimental approach, both the power model and the multi-objective optimization formulations would still be valid.

Once the model proposed in section IV for both Intel Xeon and AMD Opteron servers have been validated, we have proceeded with the analysis of results. The considered performance parameters are the temperature of both CPU and memory, as well as the frequency and voltage of the DVFS modes available to the CPU in each physical machine. These

TABLE II
CONSTANTS OBTAINED FOR POWER CURVE FITTING

Server	A	B_1	B_2	C_1	C_2	D	E	F	G_1	G_2	H_1	H_2	Error	Temp. range
Intel	14.3505	0.1110	-	-0.0011	-	0.3347	-40700	64.9494	275.702	-	-0.4644	-	11.28%	293-309K
AMD	11.2390	1.9857	-6.1703	-0.0002	0.0132	426.51	-5.3506	25.1461	-444.480	464.076	0.6977	-0.7636	9.12%	293-312K

variables modify independently the dynamic and static consumption of servers in each architecture, so different behaviors for Intel and AMD have been found. Table III shows the set of tasks used for the optimization.

TABLE III Profile of tasks allocated

Task Id	Type	# Ins	Period	Deadline	# Jobs
0	REAL	7740796	114.20	0.021	131
1	CTRL	5594832	114.21	0.015	115
2	REAL	4138643	137.12	0.011	112
3	CTRL	98156923	124.66	0.267	95
4	REAL	739437676	124.76	2.01	118
5	SOFT	2591877	124.86	0.007	103
6	SOFT	3093531	124.85	0.008	112
7	SOFT	5447445	105.76	0.015	115
8	SOFT	5722568	152.21	0.016	99

These tasks have been adapted from the TUDelft workloads archive³. The task set consists of a number of deadline-critical tasks $\tau_{hrt} = \{\tau_0, \tau_2, \tau_4\}$, a number of QoC-critical control tasks $\tau_c = \{\tau_1, \tau_3\}$, and a number soft real-time tasks $\tau_{srt} = \{\tau_5, \tau_6, \tau_7, \tau_8\}$. We assume that all tasks are independent from each other. However, due to the interference from other tasks, each task τ_i experiences a response time or delay R_i . Periods and deadlines are given in seconds. Each real-time tasks τ_{hrt} is bounded to one single host. Only control τ_c and soft tasks τ_{srt} are allowed to loss their deadline, increasing the λ parameter in the multi-objective function. Control tasks are configured with S=2.

NSGA-II has been executed with the minimum frequency in all the CPUs (labeled in the results as DVFS-MIN), the maximum frequency (labeled as DVFS-MAX) and a range of 5 possible DVFS modes (from 1 to 5). This algorithm has been compared with a more traditional approach, the EDF-VD algorithm. The overall goal is to design a priority assignment technique with the following objectives:

- All the real-time tasks au_{hrt} meet their deadlines D_{hrt} in the WCET
- The overall QoC of all the control tasks τ_c and QoS of all the soft real-time tasks τ_{srt} is maximized.
- The overall energy is minimized.

Figure 3 depicts the three obtained Pareto fronts for the Intel architecture. Both objectives have been normalized to the worst value in all the Intel and AMD optimizations (1 unit of energy = 95.6 KJ). As can be seen, the DVFS-MAX Intel framework is able to allocate all the tasks in Table III without penalizations (labeled as full feasibility). Using DVFS-MIN, the algorithm was not able to allocate all the required tasks, having to break some soft timing constraints (labeled as partial feasibility). As can be seen, there is at least one DVFS-VAR



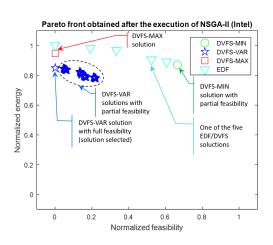


Fig. 3. Pareto front obtained with NSGA-II after optimizing the allocation of tasks over the Intel architecture.

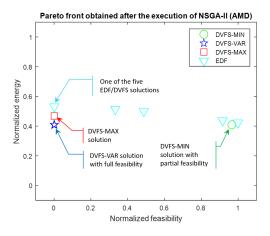


Fig. 4. Pareto front obtained with NSGA-II after optimizing the allocation of tasks over the AMD architecture.

configuration able to execute all the tasks without penalization and with less energy than DVFS-MAX and close to DVFS-MIN. Table IV shows the DVFS modes selected by the DVFS-VAR solution with full feasibility.

Similarly, Figure 4 shows the three obtained non-dominated fronts for the AMD architecture. As with the Intel scenario, the algorithm was not able to execute all the REAL, CTRL and SOFT tasks without penalization using the minimum DVFS mode (DVFS-MIN), although all the REAL tasks were properly executed. However, we found a completely feasible solution in DVFS-VAR (feasibility=0), consuming less energy than DVFS-MAX. Table IV shows the DVFS models selected by the multi-objective algorithm in the DVFS-VAR AMD optimization.

EDF was able to schedule all the tasks in both cases, but using the maximum DVFS mode and thus consuming more energy than the proposed algorithm.

TABLE IV
DVFS MODES OBTAINED BY NSGA-II PARETO FRONT IN THE
DVFS-VAR OPTIMIZATION

Platform	CPU 1	CPU 2	CPU 3	CPU 4	CPU 5	CPU 6
Intel	2	2	2	2	2	5
AMD	5	1	5			

As a result, the best DVFS configuration that can execute all the demanded services given in Table III has been found without penalizations, obtaining a high diversity in terms of energy consumption.

VII. CONCLUSIONS

CPS and Mobile Cloud Computing have collided with the lack of accurate power models for the energy-efficient provisioning of their devised infrastructures, and the real-time management of the computing facilities. In this paper, we have presented a reservation-based scheme aiming to jointly schedule deadline-critical, QoS non-critical, and QoC tasks. The work proposed in this paper has made substantial contributions in the area of power modeling of high-performance servers for Cloud computing services under timing constraints, which is an interesting and big challenge.

We have proposed an accurate power model in data centers for time constrained servers in Cloud computing, which does not only consider the workload assigned to the processing element, but also incorporates the need of considering the static power consumption and its dependency with temperature.

The proposed model has been used in a multi-objective optimization environment in which the DVFS and workload assignment have been efficiently optimized in a realistic scenario composed of Fujitsu RX300 S6 servers based on an Intel Xeon E5620 and SunFire V20z Dual Core AMD Opteron 270. Results show that the proposed multi-objective optimization framework is able to find the best DVFS configuration that can execute all the given demanded services without penalizations. In addition, the set of non-dominated solutions found presents a high diversity in terms of energy consumption.

The obtained results open a motivating research line that could enable intensely sought Green computing paradigm with hard timing constraints. Future work envisages to extend the scheduling model to integrate the concept of criticality levels.

ACKNOWLEDGMENT

This work is supported by the Spanish Ministry of Economy and Competitivity under research grants TIN2013-40968-P and TIN2014-54806-R.

REFERENCES

- [1] E. A. Lee, "CPS Foundations," in *Proceedings of the 47th Design Automation Conference*, ser. DAC '10. New York, NY, USA: ACM, 2010, pp. 737–742.
- [2] A. Banerjee, K. K. Venkatasubramanian, T. Mukherjee, and S. K. S. Gupta, "Ensuring safety, security, and sustainability of mission-critical cyber-physical systems." *Proceedings of the IEEE*, vol. 100, no. 1, pp. 283–299, 2012. [Online]. Available: http://dblp.uni-trier.de/db/journals/pieee/pieee100.html

- [3] A. Rowe, K. Lakshmanan, H. Zhu, and R. Rajkumar, "Rate-harmonized scheduling for saving energy," in *Proceedings of the 2008 Real-Time Systems Symposium*, ser. RTSS '08. Washington, DC, USA: IEEE Computer Society, 2008, pp. 113–122. [Online]. Available: http://dx.doi.org/10.1109/RTSS.2008.50
- [4] J. Chen and C. Kuo, "Energy-efficient scheduling for real-time systems on dynamic voltage scaling (DVS) platforms," in 13th IEEE International Conference on Embedded and Real-Time Computing Systems and Applications (RTCSA 2007), 21-24 August 2007, Daegu, Korea, 2007, pp. 28–38.
- [5] J. Kim, H. Kim, K. Lakshmanan, and R. R. Rajkumar, "Parallel scheduling for cyber-physical systems: Analysis and case study on a self-driving car," in *Proceedings of the ACM/IEEE 4th International Conference on Cyber-Physical Systems*, ser. ICCPS '13. New York, NY, USA: ACM, 2013, pp. 31–40. [Online]. Available: http://doi.acm.org/10.1145/2502524.2502530
- [6] P. Huang, P. Kumar, G. Giannopoulou, and L. Thiele, "Energy efficient dvfs scheduling for mixed-criticality systems," in *Proceedings of the* 14th International Conference on Embedded Software, ser. EMSOFT '14. New York, NY, USA: ACM, 2014, pp. 11:1–11:10. [Online]. Available: http://doi.acm.org/10.1145/2656045.2656057
- [7] S. K. Baruah, V. Bonifaci, G. D'Angelo, H. Li, A. Marchetti-Spaccamela, S. van der Ster, and L. Stougie, "The preemptive uniprocessor scheduling of mixed-criticality implicit-deadline sporadic task systems." in *ECRTS*, R. Davis, Ed. IEEE Computer Society, 2012, pp. 145–154.
- [8] L. Wang, G. von Laszewski, J. Dayal, and F. Wang, "Towards energy aware scheduling for precedence constrained parallel tasks in a cluster with DVFS," in Cluster, Cloud and Grid Computing (CCGrid), 2010 10th IEEE/ACM International Conference on, May 2010, pp. 368–377.
- [9] A. Beloglazov, J. Abawajy, and R. Buyya, "Energy-aware resource allocation heuristics for efficient management of data centers for cloud computing," *Future Gener. Comput. Syst.*, vol. 28, no. 5, pp. 755–768, May 2012. [Online]. Available: http://dx.doi.org/10.1016/j.future.2011. 04.017
- [10] C. Wu, R. Chang, and H. Chan, "A green energy-efficient scheduling algorithm using the DVFS technique for cloud datacenters," *Future Generation Comp. Syst.*, vol. 37, pp. 141–147, 2014.
- [11] S. Park, J. Kim, and G. Fox, "Effective real-time scheduling algorithm for cyber physical systems society," *Future Generation Comp. Syst.*, vol. 32, pp. 253–259, 2014.
- [12] C. Mobius, W. Dargie, and A. Schill, "Power consumption estimation models for processors, virtual machines, and servers," *Parallel and Distributed Systems, IEEE Transactions on*, vol. 25, no. 6, pp. 1600– 1614. June 2014.
- [13] M. M. Rafique and et al., "Power management for heterogeneous clusters: An experimental study," in *IGCC*, Washington, DC, USA, 2011, pp. 1–8.
- [14] R. I. Davis and A. Burns, "A survey of hard real-time scheduling for multiprocessor systems," ACM Comput. Surv., vol. 43, no. 4, pp. 35:1–35:44, Oct. 2011. [Online]. Available: http://doi.acm.org/10.1145/ 1978802.1978814
- [15] R. Schneider, D. Goswami, A. Masrur, M. Becker, and S. Chakraborty, "Multi-layered scheduling of mixed-criticality cyber-physical systems," *Journal of Systems Architecture (JSA)*, vol. 59, no. 10-D, 2013.
- [16] L. Abeni, N. Manica, and L. Palopoli, "Efficient and robust probabilistic guarantees for real-time tasks," J. Syst. Softw., vol. 85, no. 5, pp. 1147–1156, May 2012. [Online]. Available: http://dx.doi.org/10.1016/j.jss.2011.12.042
- [17] T. Chantem, X. Hu, and M. Lemmon, "Generalized elastic scheduling for real-time tasks," *Computers, IEEE Transactions on*, vol. 58, no. 4, pp. 480–495, 2009.
- [18] S. Narendra and A. Chandrakasan, Leakage in Nanometer CMOS Technologies, ser. Integrated Circuits and Systems. Springer, 2010.
- [19] J. Rabaey, Low Power Design Essentials, ser. Engineering (Springer-11647). Springer, 2009.
- [20] K. Deb, A. Pratap, S. Agarwal, and T. Meyarivan, "A Fast and Elitist Multiobjective Genetic Algorithm: NSGA-II," *IEEE Transactions on Evolutionary Computation*, vol. 6, no. 2, pp. 182–197, 2002.
- [21] A. Sayyad and H. Ammar, "Pareto-optimal search-based software engineering (posbse): A literature survey," in Realizing Artificial Intelligence Synergies in Software Engineering (RAISE), 2013 2nd International Workshop on, May 2013, pp. 21–27.